

9

Ray Tracing

A transportation network for light


R. Mukundan (mukundan@canterbury.ac.nz)

Department of Computer Science and Software Engineering
University of Canterbury, New Zealand.




Local Illumination Model

- A local illumination model considers only light travelling from a light source to a surface and then reflected off the surface to the eye.
 - Requires only the light source coordinates, local surface geometry and the material characteristics at a vertex.
 - Suitable for the hardware pipeline (OpenGL, Direct3D)
- Does not consider occlusions and transmittance of light.



Global Illumination


- The illumination at a given point is a combination of the light received from a source and the light reflected from other surfaces in the scene.
- Considers the effects of occlusions, surface reflections, light transmission through a medium (direct transmittance and refractions), and indirect illumination.



Global Illumination Methods

Radiosity:


- A scene illumination can be considered as an equilibrium state for radiant energy transfers between surface elements.
- Gives good results for diffuse illumination, but specularity is not handled.
 - Useful for modelling area light sources, colour bleeding, soft shadows.



Cornell Box


Ray Tracing (Backward Ray Tracing)


- Traces ray outwards from the eye to objects and light sources. The opposite of reality.
- Use secondary rays to determine shadows and to model mirror reflection and refraction
- We can easily generate many global illumination effects
- But:
 - Doesn't handle diffuse inter-reflections
 - e.g. “colour bleed” from a bright red wall to an adjacent white wall
 - Computationally expensive



“Tracing” a ray

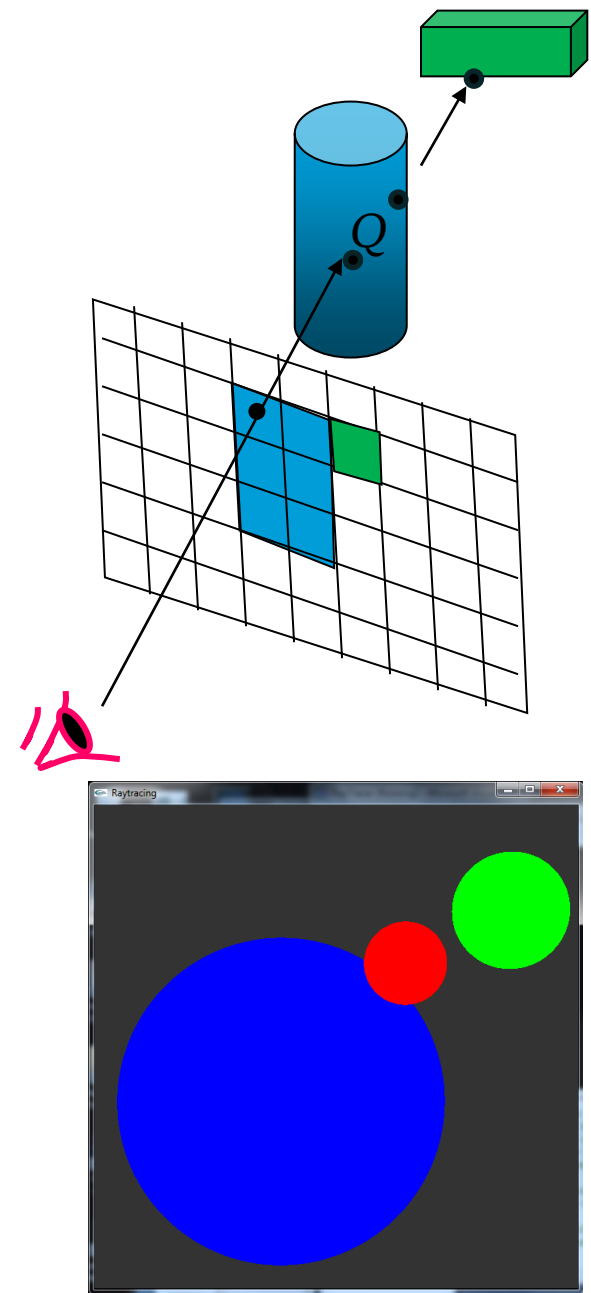
- Compare the ray with all scene objects and compute the closest point of intersection Q , and obtain the intersecting object's index.
- Compute the colour value at the point of intersection Q .
- Generate a shadow ray to determine if the point Q is in shadow (2)
- If the surface is reflective, recursively trace the reflected ray at Q (3)
- If the surface is refractive, recursively trace the refracted ray at Q (4)
- Add the colour contributions from (2), (3) and (4), and return the colour value.





Ray Casting

- Ray tracing without secondary rays.
- Trace a ray from the view point (called the primary ray) through each “pixel” on the image plane
 - Test each surface to determine if it is intersected by the ray.
 - Compute the points of intersection on each primary ray.
 - Get the point of intersection Q that is closest to the eye.
 - Use the colour of the object on which Q lies as pixel colour.



Ray Casting + Phong Lighting

- At the point of intersection Q , compute the colour value using an illumination model.
- Simplified Phong Illumination:

Assumptions:

Light's ambient color $A = (0.2, 0.2, 0.2)$

Light's diffuse and specular color $= (1, 1, 1)$

M = Material Color (ambient and diffuse)



Material's specular color $= (1, 1, 1)$

$$\text{Col} = \underset{\substack{\uparrow \\ \text{Ambient}}}{AM} + \underset{\substack{\uparrow \\ \text{Diffuse}}}{M(l \cdot n)} + \underset{\substack{\uparrow \\ \text{Specular}}}{(1, 1, 1)(r \cdot v)^f}$$


Ambient

Diffuse

Specular



Shadow Ray




- Trace a ray from the point of intersection Q towards the light source L (2)
- If the shadow ray hits an object, and if the point of intersection R is between the light source and the object A (i.e., $RQ < LQ$), then the point Q is in shadow.

If Q is in shadow

$$I_A = AM$$

Else



$$I_A = AM + M (I \cdot n) + (1, 1, 1) (r \cdot v)^f$$



Reflections

- If the surface is reflective, then a secondary ray along the direction of reflection is traced.
- If this secondary ray meets a surface at a point with intensity I_B , then $\rho_r I_B$ is added to the pixel color
 - ρ_r is a scale factor (< 1), called the coefficient of reflection
 - ρ_r represents how much of colour I_B is reflected on the surface A .
- The colour of the pixel is now

$$I = I_A + \rho_r I_B$$




Refractions

- If the surface is transparent, a secondary ray along the direction of refraction is traced.
- If this secondary ray meets a surface at a point with intensity I_C , then $\rho_t I_C$ is added to the pixel color.
 - ρ_t is a scale factor (<1), called the coeff. of transmission



- The colour of the pixel is now

$$I = I_A + \rho_t I_C$$

(Compare with a similar equation on Slide 12)



Binary Ray-tracing Tree




- • Left branches represent reflections
- • Right branches represent transmission paths

What is a “ray”?

A ray is specified using

- A point (the source of the ray): $p_0 = (x_0, y_0, z_0)$
- A vector (the direction of the ray): $d = (l, m, n)$
- The vector must be converted to a unit vector.

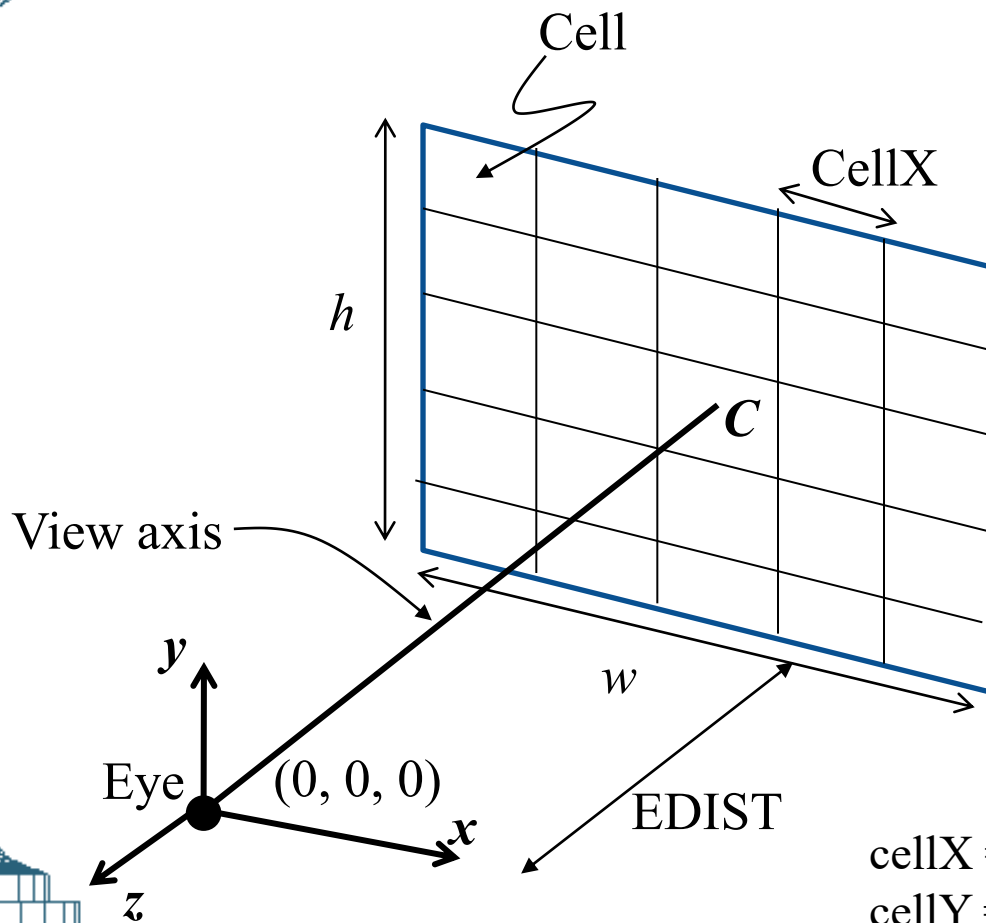
Any point on the ray can be represented using a single parameter t . The value of t denotes the distance from the source to that point.



Ray's Equation

```
glm::vec3 p0(x0, y0, z0);  
glm::vec3 dir(l, m, n);  
Ray ray = Ray(p0, dir);  
ray.normalize();
```

Ray Tracing Setup



w = Width of screen in world units (eg. 5 units).

h = Height of screen in world units (eg. 5 units).

EDIST = Eye dist in world units (eg. 10 units)

$XMIN = -w/2$; $XMAX = w/2$;

$YMIN = -h/2$; $YMAX = h/2$;


NUMDIV = Number of subdivisions along x, y directions.

cellX = cell width = $(XMAX - XMIN) / NUMDIV$

cellY = cell height = $(YMAX - YMIN) / NUMDIV$


cell indices: $i, j = 0, \dots, NUMDIV - 1$.

Primary Ray



- Ray position: $(0, 0, 0)$
- Ray direction d : $((XMIN+(i+0.5)*CellX, (YMIN+(i+0.5)*CellY, -EDIST))$
- Normalize the above direction.


Secondary Rays



Shadow ray: A ray that originates at a point of intersection with an object, and directed towards a light source.

Position = q (point of intersection)

Direction = $L - q$ normalized.



Reflection ray: A ray from the intersection point towards the direction of reflection of the incident ray (Used only for reflective surfaces)

Position = q

Direction:

$$r = -2(d \cdot n)n + d \text{ normalized.}$$

(We will use a GLM function to compute r)

GLM Functions for Ray Tracing


- Reflection:

$\mathbf{r} = \text{glm::reflect}(\mathbf{d}, \mathbf{n});$

\mathbf{d} : Unit incident vector

\mathbf{n} : Unit normal vector

The reflection vector also is a unit vector.



- Refraction:


$\mathbf{g} = \text{glm::refract}(\mathbf{d}, \mathbf{n}, \text{eta});$

\mathbf{d} : Unit incident vector

\mathbf{n} : Unit normal vector

The refraction vector also is a unit vector.

eta = Ratio of refractive indices = η_1/η_2



Reflections


- Computation of specular reflections:

$\mathbf{r} = \text{glm::reflect}(-\mathbf{l}, \mathbf{n});$

\mathbf{l} : Unit light source vector


Specular term = $(\mathbf{r} \cdot \mathbf{v})^f$.

\mathbf{v} : View vector = $-\mathbf{d}$



- Surface reflections:


$\mathbf{r} = \text{glm::reflect}(\mathbf{d}, \mathbf{n});$




```
if(ray.xindex == 0 && step < MAX_STEPS)
{
    glm::vec3 refl = glm::reflect(ray.dir, normalVector);
    Ray reflectedRay(ray.xpt, refl);
    glm::vec3 reflCol = trace(reflectedRay, step+1);
    colorSum = colorSum + (0.8f*reflCol);
}
```

Index of Refraction

- Light travels at speed c/η in a medium with index of refraction η .
- Common values of index of refraction:
 - Air 1.
 - Water 1.33
 - Glass 1.5
 - Diamond 2.4
- Snell's Law of Refraction:
 - $\eta_1 \sin \theta_1 = \eta_2 \sin \theta_2$




Refracted Ray





$$\begin{aligned} u \sin \theta_i - n \cos \theta_i &= d \\ u \sin \theta_t - n \cos \theta_t &= g \end{aligned}$$

$$\mathbf{g} = \left(\frac{\eta_1}{\eta_2} \right) \mathbf{d} - \left(\frac{\eta_1}{\eta_2} (\mathbf{d} \cdot \mathbf{n}) + \cos \theta_t \right) \mathbf{n}$$

$$\cos \theta_t = \sqrt{\left(1 - \left(\frac{\eta_1}{\eta_2} \right)^2 (1 - (\mathbf{d} \cdot \mathbf{n})^2) \right)}$$



Sphere Refraction



```
 $n$  = sceneObjects[ray.xindex]->normal(ray.xpt);
```

```
 $g$  = glm::refract( $d$ ,  $n$ , eta);
```


```
Ray refrRay1(ray.xpt,  $g$ )
```

```
refrRay1.closestPt(sceneObjects);
```

```
 $m$  = sceneObjects[refrRay1.xindex]->normal(refrRay1.xpt);
```

```
 $h$  = glm::refract( $g$ ,  $-m$ , 1.0f/eta);
```


Sphere Refraction




$$\text{eta} = 1/1.5$$

Sphere Refractions



$$\text{eta} = 1/1.01$$



$$\text{eta} = 1/1.003$$

Shadows



Bad



Good



Better

Ray-plane intersection

- Given a point a on a plane, and the normal vector n of the plane, the plane's equation can be written as

$$(p - a) \bullet n = 0$$

- A ray is given by the equation


$$p = p_0 + t d.$$

- At the point of intersection, both equations are true.

Therefore, $(p_0 + t d - a) \bullet n = 0$

- From the above equation, we get

$$t = \frac{(a - p_0) \bullet n}{d \bullet n}$$



Ray-sphere intersection

- Equation of a sphere centred at C with radius r is
$$(p - C) \bullet (p - C) = r^2$$
- Consider a ray given by the equation
$$p = p_0 + t d.$$
- At the point of intersection, both equations are true.


Therefore, $(p_0 + t d - C) \bullet (p_0 + t d - C) = r^2$

$$(s + t d) \bullet (s + t d) = r^2, \quad \text{where } s = p_0 - C$$

$$(d \bullet d) t^2 + 2 (s \bullet d) t + (s \bullet s) - r^2 = 0.$$


Since d is a unit vector, we get

$$t = -(s \bullet d) \pm \sqrt{(s \bullet d)^2 - (s \bullet s) + r^2}$$



Ray intersection with polygonal objects


- Use the plane equation of each triangle to get the point of intersection with the ray.
- Check if the point of intersection lies within the triangle.



Intersection :

$$t = \frac{(A - p_0) \cdot n}{d \cdot n}$$

$$q = p_0 + t d$$




Point inclusion test :

If the cross products $(u_A \times v_A)$, $(u_B \times v_B)$, $(u_C \times v_C)$ have the same sign, then the point q is inside the triangle.


Ray intersection with polygonal objects

- Complex polygonal objects will require a large amount of ray-triangle intersection tests.
- Bounding volume hierarchies and spatial subdivision methods (*kd*-Trees, Octrees) are used to reduce the number of ray-primitive intersection tests.



Texture Mapping

- 2D texturing:
 - Similar to OpenGL texturing, but does not use texture coordinates or texture memory.
 - Map the coordinates of the point of intersection (x, y, z) to image coordinates, and assign the colour of the pixel to that point.



- Procedural texturing:
 - Define functions to map (x, y, z) coordinates to (r, g, b) colour values. $r = r(x, y, z)$; $g = g(x, y, z)$; $b = b(x, y, z)$;

Texture Mapping

```
...
#include "Ray.h"
#include "TextureBMP.h"
#include <GL/glut.h>
using namespace std;
...
TextureBMP texture;
...
void initialize()
{
    ...
    texture = TextureBMP("myPicture.bmp");
}
...
glm::vec3 trace(Ray ray, int step)
{
    ...
    col = texture.getColorAt(s, t);
}
```

Lab08


TextureBMP.h
TextureBMP.cpp

Note:

$0 \leq s \leq 1$

$0 \leq t \leq 1$


Texturing a Plane




```
if(ray.xindex == 3)
{
    texcoords = (ray.xpt.x - a1)/(a2-a1);
    texcoordt = (ray.xpt.y - b1)/(b2-b1);
    col = texture.getColorAt(texcoords, texcoordt);
}
```

Anti-Aliasing


The ray tracing algorithm samples the light field using a finite set of rays generated through a discretized image space. This results in distortion artefacts such as jaggedness along edges of polygons and shadows.



"Normal" Ray Tracing




- Supersampling: Generate several rays through each square pixel (eg. divide the pixel into four equal segments) and compute the average of the colour values.



Anti-Aliasing

- Adaptive Sampling: As shown on the previous slide, each pixel is divided into four “sub-pixels”. Primary rays are generated through the centres of each sub-pixel. If the colour value along any ray varies significantly from the other three, that sub-pixel is split further into four sub-pixels, and more rays are generated through them.




Cylinder

- A cylinder at the origin with axis along the y -axis, radius R and height h is given by

$$x^2 + z^2 = R^2$$

$$0 \leq y \leq h$$

- Normal vector at (x, y, z)
(un-normalized) $n = (x, 0, z)$
Normalized $n = (x/R, 0, z/R)$



Ray – Cylinder Intersection

- A cylinder at (x_c, y_c, z_c) , with axis parallel to the y -axis, radius R and height h is given by

$$(x - x_c)^2 + (z - z_c)^2 = R^2$$

$$0 \leq (y - y_c) \leq h$$

- Normal vector at (x, y, z)

(un-normalized) $\mathbf{n} = (x - x_c, 0, z - z_c)$


(normalized) $\mathbf{n} = ((x - x_c)/R, 0, (z - z_c)/R)$

Ray equation:

$$x = x_0 + d_x t; \quad y = y_0 + d_y t; \quad z = z_0 + d_z t;$$

- Intersection equation:

$$t^2 (d_x^2 + d_z^2) + 2t \{ d_x (x_0 - x_c) + d_z (z_0 - z_c) \} + \{ (x_0 - x_c)^2 + (z_0 - z_c)^2 - R^2 \} = 0.$$



Cone

- Consider a cone with centre of base at (x_c, y_c, z_c) , axis parallel to the y -axis, radius R , and height h :

- Important equations:

$$\tan(\theta) = R/h \quad (\theta = \text{half cone angle})$$

For any point (x, y, z) on the cone,



$$(x - x_c)^2 + (z - z_c)^2 = r^2$$

where,
$$r = \left(\frac{R}{h}\right)(h - y + y_c)$$

- Surface normal vector (normalized):

$$\mathbf{n} = (\sin \alpha \cos \theta, \sin \theta, \cos \alpha \cos \theta)$$

where
$$\alpha = \tan^{-1} \left(\frac{x - x_c}{z - z_c} \right)$$



Ray-Cone Intersection

- Equation of a cone with base at (x_c, y_c, z_c) , axis parallel to the y -axis, radius R , and height h :

$$(x - x_c)^2 + (z - z_c)^2 = \left(\frac{R}{h}\right)^2 (h - y + y_c)^2$$

- Ray equation:

$$x = x_0 + d_x t; \quad y = y_0 + d_y t; \quad z = z_0 + d_z t;$$

- The points of intersection are obtained by substituting the ray equation in the cone's equation and solving for t .

